

# CSE6331: Cloud Computing

Leonidas Fegaras


University of Texas at Arlington  
©2016 by Leonidas Fegaras

Graph Processing with Pregel and Giraph

# Vertex-Centric Distributed Graph Processing


- Google's Pregel
  - Distributed system for large scale graph processing
  - Vertex-centric graph processing
  - Focuses on independent local actions on graph vertices
  - Vertex-centric = “think like a vertex”
  - Fault-tolerant
  - Bulk Synchronous Parallel (BSP) as execution model
- Apache Giraph is an open source implementation of Pregel
  - Part of Hadoop ecosystem
  - In-memory distributed execution based on BSP
  - Fault-tolerance by checkpointing
- GraphX
  - Spark's API for graphs and graph-parallel computation

# The Bulk Synchronous Parallel (BSP) Execution Model



# Vertex-Centric Computations

- Each vertex has an id, a vertex value, a list of edges to adjacent neighbors
- Each edge has an id (the destination id) and an edge value
- Each vertex used at each superstep can recompute its vertex value and send messages to other vertices, which are delivered before the next superstep
- Additional features: termination votes, combiners, aggregators, topology mutations



# Model of Computation

- Computation proceeds in supersteps
  - In a superstep, vertices execute the same user-defined function *compute* in parallel
- In a superstep, a vertex can:
  - Modify its vertex value
  - Modify the value of its outgoing edges
  - Receive messages sent to it in the previous superstep
  - Send messages to other vertices to be received in the next superstep
  - Add/remove outgoing edges
- Edges may have value, but have no associated computation
- This is a pure message-passing model: no remote reads or shared memory

# PageRank in Pregel

- $P_i$ : PageRank of node  $i$
- $C_i$ : number of outgoing links from node  $i$

$$P_i = \sum_{j \rightarrow i} \frac{P_j}{C_j}$$

```
class PageRankVertex {  
    void compute ( Iterator messages ) {  
        if (getSuperstep() > 0)  
            setVertexValue(sum(messages));  
        if (getSuperstep() < k)  
            sendMessageToAllNeighbors(pageRank / getNumOutEdges());  
        else voteToHalt();  
    }  
}
```

# The Giraph API

- Need to implement the `Vertex<I,V,E,M>` interface
  - I - Vertex id (`WritableComparable`)
  - V - Vertex data (`Writable`)
  - E - Edge data (`Writable`)
  - M - Message data (`Writable`)

<b>void</b> addEdge( <code>Edge&lt;I,E&gt;</code> edge)	Add an edge to this vertex
<code>Iterable &lt;Edge&lt;I,E&gt;&gt;</code> getEdges()	Get the out-edges of this vertex
<code>E</code> getEdgeValue( <code>I</code> vid)	Get the edge value with the given target vertex id
<code>I</code> getId()	Get the vertex id
<code>V</code> getValue()	Get the vertex value
<b>void</b> removeEdges( <code>I</code> vid)	Remove all edges pointing to the given vertex id
<b>void</b> setEdges ( <code>Iterable &lt;Edge&lt;I,E&gt;&gt;</code> edges)	Set the outgoing edges for this vertex
<b>void</b> setEdgeValue( <code>I</code> vid , <code>E</code> value)	Set the edge value of the edge to the target vid
<b>void</b> setValue( <code>V</code> value)	Set the vertex data
<b>void</b> voteToHalt()	Halt until a message is sent to this vertex
<b>void</b> sendMessage( <code>I</code> vid , <code>M</code> m)	Send a message <code>m</code> to the node <code>vid</code>
<b>void</b> sendMsgToAllEdges( <code>M</code> m)	Send a message <code>m</code> to all outgoing edges

# Connected Components in Giraph

```
public void compute ( Iterator <IntWritable> messages ) {  
    int m = getValue().get();  
    for ( IntWritable message : messages )  
        m = Math.min(m,messages.next().get());  
    if (getSuperstep() == 0 || m != getValue().get())  
        setValue(new IntWritable(m));  
    sendMsgToAllEdges(m);  
    voteToHalt();  
}
```

Every vertex votes to halt on every superstep, but

- 1 if there is at least one vertex whose value changed, it will send messages to neighbors
- 2 if a vertex receives a message, it will wake up and continue, although it has voted to halt



# Single Source Shortest Paths in Giraph

- 1 find the minimum value arriving on any message
- 2 if that value is less than the current value of the vertex
- 3 the minimum will be adopted as the vertex value
- 4 the value plus the edge value will be sent along every outgoing edge

```
public void compute ( Iterable<DoubleWritable> messages ) {  
    double minDist = Double.MAX_VALUE;  
    for ( DoubleWritable message : messages )  
        minDist = Math.min(minDist, message.get());  
    if (minDist < getValue().get()) {  
        setValue(new DoubleWritable(minDist));  
        for ( Edge<LongWritable, FloatWritable> edge: getEdges() ) {  
            double distance = minDist + edge.getValue().get();  
            sendMessage(edge.getTargetVertexId(), new DoubleWritable(distance));  
        }  
    }  
    voteToHalt();  
}
```